Out of Reach: Quantifying Housing Affordability across Dutch Provinces (2013-2025)

Tutorial 5 group 8 - Noureddine El Kaouakib (2863711), Imad el Fakiri (2863975), Achraf Saih (2865196

2026-06-13

Contents

1 Identification of the social problem

1.1 Describe the social problem

Home ownership in the Netherlands has moved out of reach for a growing share of households. The scale of the problem is stark: the CPB's first Housing Market Accessibility Monitor found that the share of homes an average-income household could afford fell from 61% a decade ago to just 21% in 2024, and that such a household now falls roughly EUR 100,000 short of the mortgage needed for an average home (CPB, 2026). House prices rose far faster than incomes over this period, so the number of years of income required to buy an average home climbed steeply. This matters for wealth inequality (owners gain, would-be buyers are locked out), for household debt – the Netherlands has among the highest mortgage debt in Europe (DNB, 2024) – and for regional divergence: if affordability worsens unevenly, it shapes where people can afford to live and work.

Existing analyses (CPB, 2026; DNB, 2024) establish the national trend but focus on the country as a whole or on the major cities. The gap we address is the **provincial** picture: whether affordability deteriorated evenly across the twelve provinces or concentrated in particular regions, and how it responded to the 2022 interest-rate shock. We quantify affordability as a **price-to-income ratio** per province (average purchase price divided by median household income) and track it from 2013 to 2024, using the rapid ECB rate rises of 2022 as a natural event.

2 Data sourcing

2.1 Description and limitations

We merge **two distinct public CBS datasets**, joined on province x year:

Dataset	Unit	Key variable used	Source / period
Existing own homes; average purchase prices	province x year	`avg_price` (EUR)	CBS / Kadaster, 2013-2025
Income of persons; characteristics, region	province x year	median standardised household income (EUR)	CBS, 2013-2024

Limitations (state honestly - rubric rewards this):

- The two series do not fully overlap: prices run to **2025**, income only to **2024**, and the 2024 income figure is **provisional** (CBS flags it `*`). The affordability ratio is therefore computed only for **2013-2024**; the 2025 price is shown for context but has no matching income.
- "Average purchase price" is a mean and is sensitive to the mix of homes sold; median income is corrected for household size (standardised) but the two are not measured on the same statistical unit (transactions vs persons).
- Province-level averages hide large within-province variation (e.g. Amsterdam vs rural Noord-Holland).

2.2 Loading the data

```
## Rows: 144
## Columns: 6
## $ province          <chr> "Drenthe", "Drenthe", "Drenthe", "Drenthe", "Drenth~
## $ year              <int> 2013, 2014, 2015, 2016, 2017, 2018, 2019, 2020, 202~
```

```
## $ avg_price          <dbl> 181850, 185239, 194566, 205504, 213562, 228433, 241~
## $ med_income         <dbl> 23900, 24700, 25100, 26200, 27000, 27400, 28800, 30~
## $ price_income_ratio <dbl> 7.608787, 7.499555, 7.751633, 7.843664, 7.909704, 8~
## $ ratio_index        <dbl> 100.00000, 98.56440, 101.87739, 103.08692, 103.9548~
```

3 Quantifying

3.1 Data cleaning

Cleaning steps (full code in `R/01_clean_merge.R`), each with a rationale:

1. **Prices, wide -> long.** The CBS price file stores years as 13 columns; we pivot to one row per province x year so it can be joined to income.
2. **Income, 13 files -> one table.** Each income file is a single year with the year on its "Perioden" line; we read all of them, extract the year, and stack.
3. **Drop the duplicate 2018 file.** Two supplied income files are identical 2018 exports; keying on province x year and `distinct()` collapses them.
4. **Keep the 12 provinces only.** Income files also list four cities (Amsterdam, Rotterdam, Den Haag, Utrecht gemeente); we filter to `(PV)` rows.
5. **Fix Dutch number formats.** Income uses a decimal comma and is in thousands of euros, so we replace "," with "." and multiply by 1000.
6. **Inner join** keeps only years present in both -> a complete 12 x 12 panel.

The full script is in `R/01_clean_merge.R` and in the repo. The two steps that need the most care are the wide-to-long pivot of the price file and the Dutch decimal-comma fix in the income files; both are shown below.

```r
# Prices: wide (one column per year) -> long (one row per province-year)
prices <- price_raw %>%
  rename(region = X1) %>%
  filter(str_detect(region, "\\(PV\\)$")) %>% # provinces only
  set_names(c("region", as.character(2013:2025))) %>%
  pivot_longer(-region, names_to = "year", values_to = "avg_price")

# Income: read every column as text so the Dutch decimal comma in "23,9" is not
# misread as 239, then comma -> point, to number, x1000 (values are in 1000s euros)
med_income <- as.numeric(str_replace(med_std_income_k, ",", "."))*1000

# Merge on province + year (inner join -> the 2013-2024 overlap) and build the
# two analysis variables; distinct() collapses the duplicate 2018 income file.
panel <- prices %>%
  inner_join(income, by = c("province", "year")) %>%
  group_by(province) %>%
  mutate(price_income_ratio = avg_price/med_income,
         ratio_index = 100* price_income_ratio/price_income_ratio[year ==  2013]) %>%
  ungroup()
```

3.2 Describing the variables

Table 1: Variables in the merged panel.

Variable	Type	Meaning
province	character	Dutch province
year	integer	calendar year
avg_price	numeric	average purchase price of existing homes (EUR)
med_income	numeric	median standardised household income (EUR)
price_income_ratio	numeric	avg_price / med_income (years of income to buy)
ratio_index	numeric	ratio relative to province's 2013 level (=100)

3.3 Creating new variables

Two constructed variables (rubric: >= 2 variables with analytical value):

- **`price_income_ratio`** = average price / median income. Interprets directly as the number of years of median income needed to buy an average home - the standard affordability measure.
- **`ratio_index`** = each province's ratio relative to its own 2013 value (=100). Lets us compare the *speed of deterioration* across provinces on a common scale, independent of their starting level.

Table 2: Price-to-income ratio by period.

period	mean_ratio	min_ratio	max_ratio
2013-2015	8.20	7.24	10.34
2016-2021	9.55	7.40	14.77
2022-2024	10.90	8.90	15.25

The `ratio_index` variable (2013 = 100 for every province) makes the *speed* of deterioration directly comparable across provinces, regardless of their starting affordability level:

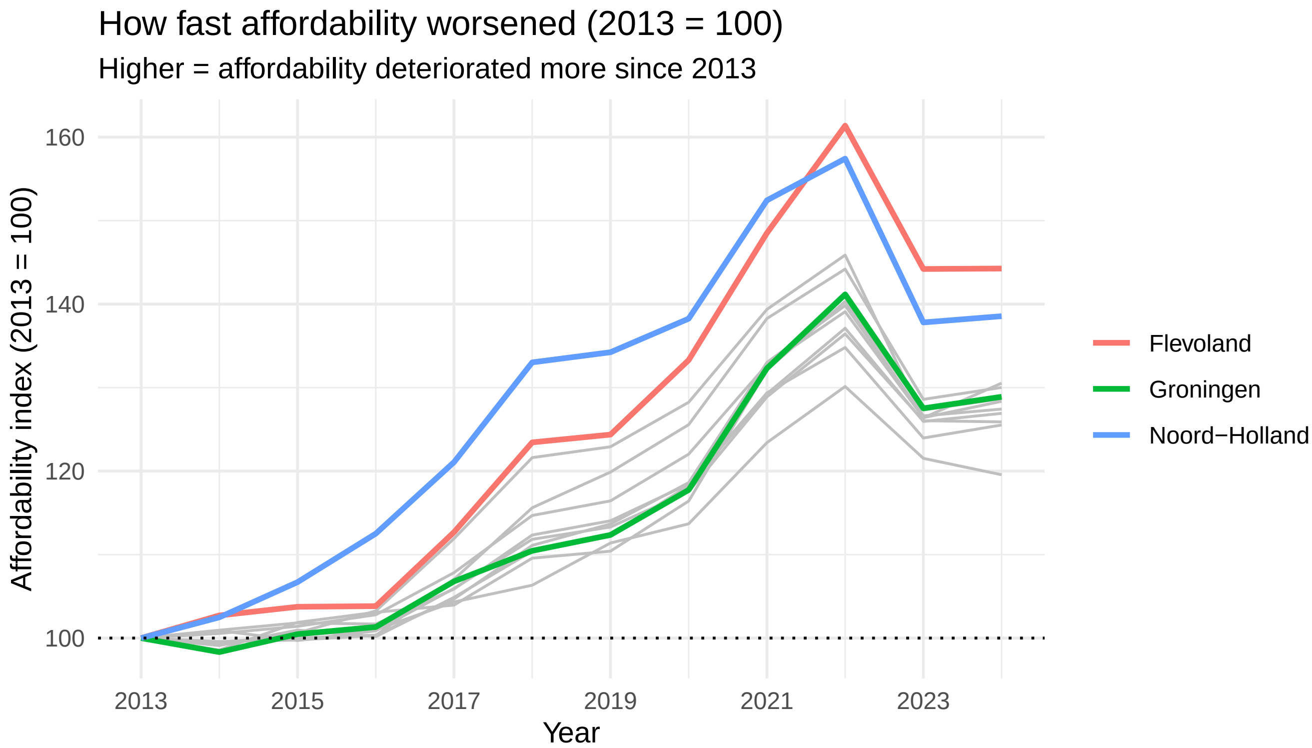

Indexing each province to its own 2013 level reveals a pattern the absolute ratios hide: the *fastest* deterioration did not occur in the most expensive province. Flevoland's affordability index reached 144 by 2024 (a 44% worsening since 2013) – the steepest of all twelve – ahead of Noord-Holland (139) and Utrecht (131). So while the Randstad has the worst absolute affordability, several cheaper provinces are catching up fast in relative terms; Zeeland, at 120, deteriorated least.

3.4 Temporal variation

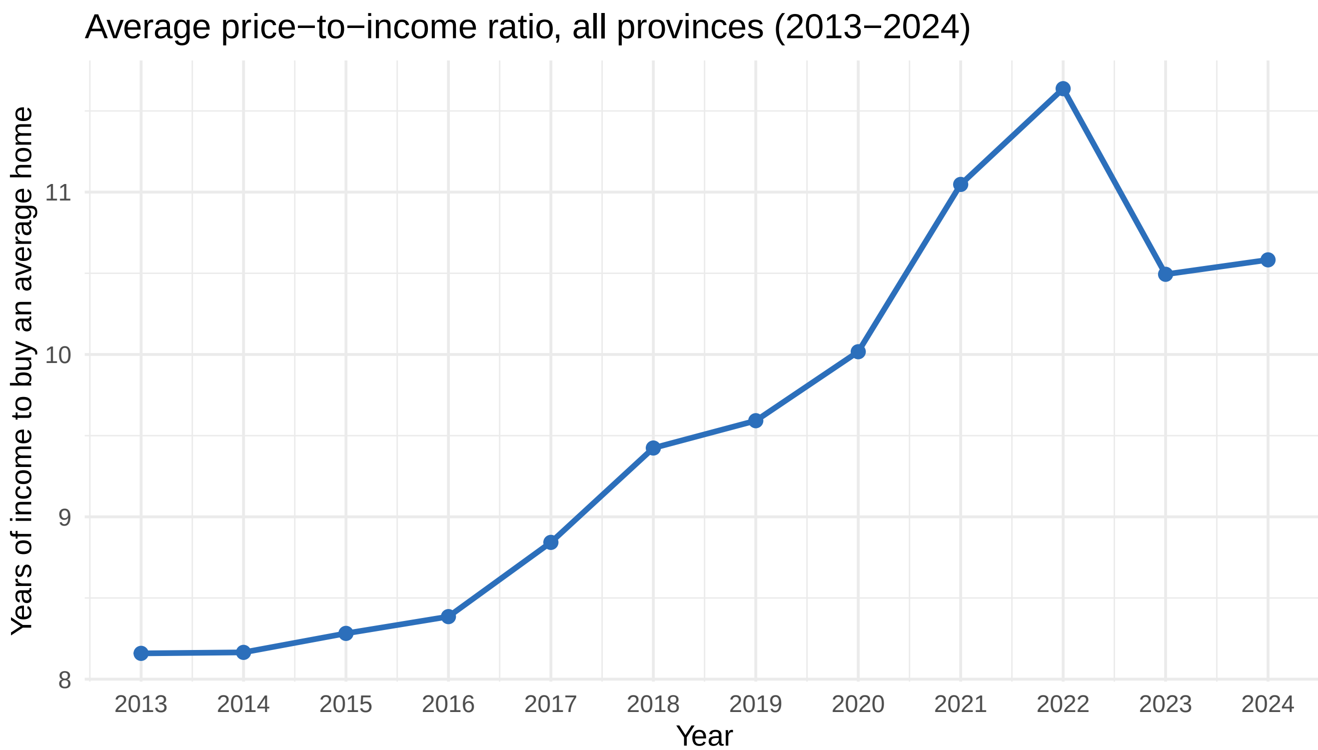

The national average ratio was essentially flat at about 8.2 years of income in 2013-2014, then climbed steadily for eight years to a peak of 11.6 in 2022 – meaning an average home went from costing roughly eight years of median income to nearly twelve. After the 2022 interest-rate shock the ratio fell to 10.5 in 2023 and stayed near that level (10.6) in 2024. Affordability therefore eased slightly after 2022 but remained far worse than at any point before 2020; the dip is a partial pause, not a recovery.

3.5 Sub-population variation

Here the "sub-populations" are the provinces (grouped/facetted comparison).

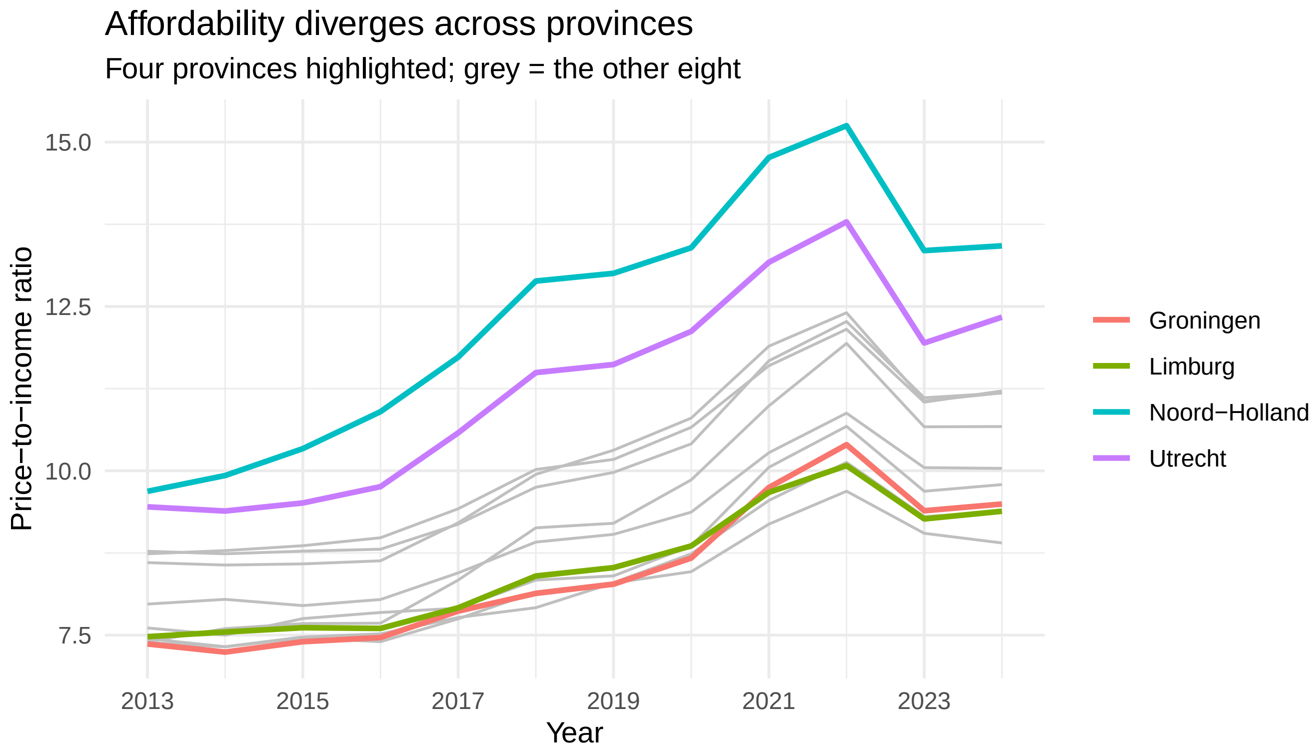

The provinces fan out over time rather than moving together. In 2013 the ratio ranged narrowly (roughly 7-10 years); by 2024 the spread had widened markedly, from Zeeland at 8.9 years to Noord-Holland at 13.4 – a gap of about 4.5 years of income between the most and least affordable province. Noord-Holland and Utrecht (the two highlighted Randstad provinces) pull away from the pack throughout, while Groningen and the northern provinces stay at the bottom. The widening gap means affordability is not just worsening nationally but becoming more regionally unequal.

3.6 Spatial variation

We map the 2024 ratio onto real province boundaries (CBS/PDOK, read with the `sf` package) - a choropleth rather than a placeholder chart.

Housing affordability by province, 2024

Average home price expressed in years of median household income

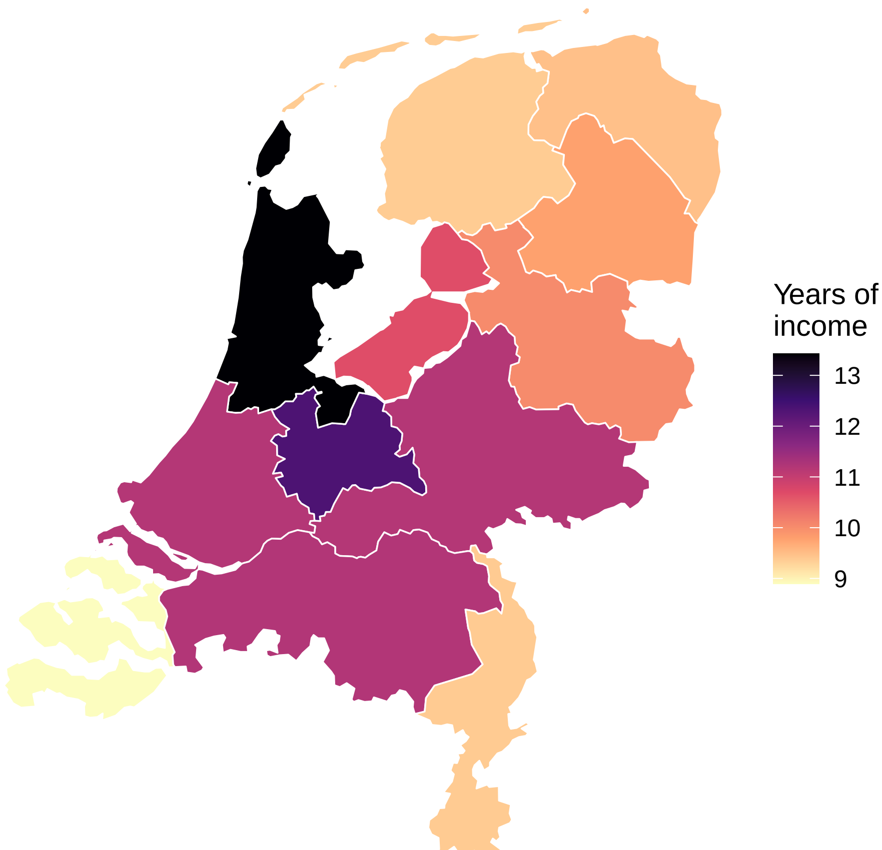

Data: CBS. Boundaries: CBS/PDOK via cartomap.

The map confirms a clear urban-core versus periphery pattern. The Randstad provinces – Noord-Holland (13.4), Utrecht (12.3) and Zuid-Holland (11.2) – form a high-ratio cluster, with Noord-Brabant and Gelderland close behind. The most affordable provinces sit at the edges of the country: Zeeland (8.9), Fryslan (9.4), Limburg (9.4) and Groningen (9.5). This mirrors the CPB's finding that accessibility problems are greatest in and around the major cities (CPB, 2026), and points to housing demand concentrated near economic centres as the spatial driver. A caveat: province-level means mask large within-province differences – Amsterdam is far less affordable than rural Noord-Holland.

3.7 Event analysis

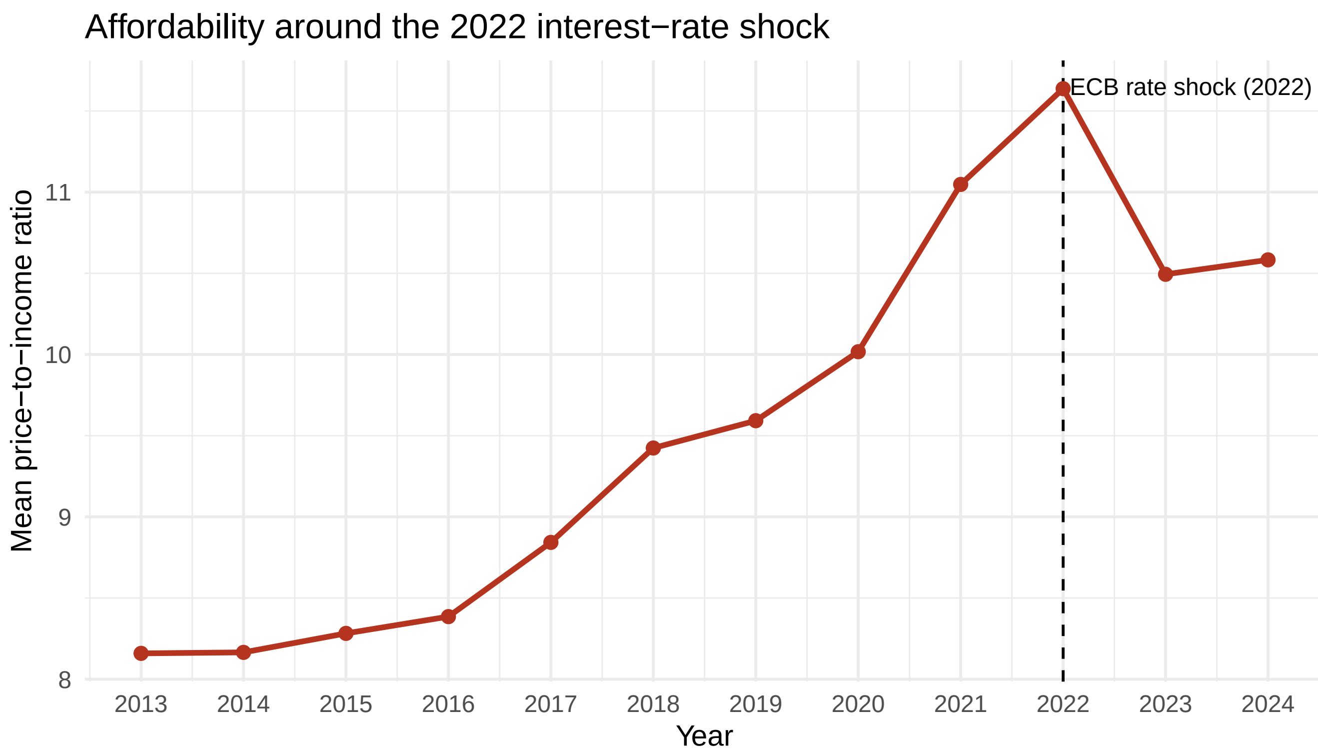

The ECB began raising rates sharply in 2022, and the dashed line marks that event. The ratio peaks exactly at 2022 (11.6) and falls for the first time in the series in 2023 (10.5) – consistent with higher mortgage rates cutting borrowing capacity and pushing prices down. **Causal caveat:** this is an association, not proof. The post-2022 easing reflects several forces at once. DNB notes that the price fall ended in summer 2023 once rates stopped climbing and incomes began rising, with prices recovering thereafter (DNB, 2024); rising incomes also lift our denominator independently of the rate shock. So the rate rise is one plausible driver of the turning point, but separating it from income growth and the structural housing shortage would require a more formal design.

4 Discussion

Our analysis quantifies how housing affordability in the Netherlands worsened between 2013 and 2024 and how it varies across provinces. Three findings stand out. First, affordability deteriorated sharply and nationally: the average home went from costing about 8.2 years of median household income in 2013 to a peak of 11.6 years in 2022. Second, the burden is regionally unequal and the gap is widening – by 2024 the ratio ranged from 8.9 in Zeeland to 13.4 in Noord-Holland, with a clear Randstad-versus-periphery divide. Third, the index analysis shows the fastest *relative* decline occurred in Flevoland rather than the most expensive province, so cheaper regions are not insulated. These results align with and add regional detail to the CPB's national finding that average earners can increasingly no longer afford to buy (CPB, 2026).

Around the 2022 interest-rate shock the ratio turned down for the first time, but we are careful not to overclaim: incomes rising, prices falling, and the structural housing shortage all coincide, and DNB documents that prices resumed rising from late 2023 (DNB, 2024). We read the 2022 turning point as an association consistent with the rate shock, not as a causal estimate.

Limitations temper these conclusions. We use average purchase prices against median standardised income, two different statistical units, so the ratio is an indicator rather than a literal "years to buy" for a typical buyer. Province-level means hide large within-province variation (notably the big cities). Income data end in 2024 (provisional) while prices run to 2025, so the ratio cannot cover the most recent year. Future work could use a real choropleth at municipal level and a proper difference-in-differences or interrupted-time-series design to test the rate-shock effect more rigorously.

5 Reproducibility

5.1 GitHub repository link

`https://github.com/noureddineek/Housing_affordability`

5.2 How this is reproduced

```
housing_project/
  README.md                     # what the project is + how to run it
  run_all.R                     # reproduces everything end to end
  housing_affordability.Rmd     # this report
  R/01_clean_merge.R            # cleaning + merge -> data/panel.csv
  R/02_make_map.R               # downloads boundaries + builds the choropleth
  data/raw/                     # the unmodified CBS downloads + province GeoJSON
  data/panel.csv                # generated
  output/                       # knitted pdf + map png
```

6 References

The data and sources used in this report:

- Centraal Bureau voor de Statistiek (CBS). (n.d.). *Bestaande koopwoningen; gemiddelde verkoopprijzen, regio* [Existing own homes; average purchase prices, region] [Data set]. CBS StatLine. https://opendata.cbs.nl/
- Centraal Bureau voor de Statistiek (CBS). (n.d.). *Inkomen van personen; persoonskenmerken, regio (indeling 2025)* [Income of persons; characteristics, region] [Data set]. CBS StatLine. https://opendata.cbs.nl/
- Centraal Planbureau (CPB). (2026). *Monitor toegankelijkheid woningmarkt* [Housing Market Accessibility Monitor]. CPB Netherlands Bureau for Economic Policy Analysis. https://www.cpb.nl/
- De Nederlandsche Bank (DNB). (2024). *Housing market.* https://www.dnb.nl/en/current-economic-issues/housing-market/
- Cartomap. (2023). *provincie_2023* [Province boundary GeoJSON, CBS/PDOK source]. https://cartomap.github.io/nl/